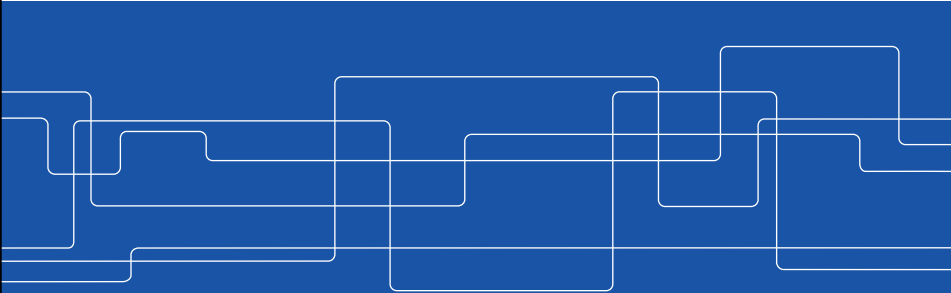


KTH ROYAL INSTITUTE  
OF TECHNOLOGY

# Scheduling Aperiodic and Sporadic Tasks in Priority-Driven Systems


Ingo Sander and Matthias Becker  
Contact: [ingo@kth.se](mailto:ingo@kth.se)

Liu: Chapter 7  
Buttazzo: Chapter 5



IL2212 – Embedded Software

1



## Outline

- System Model and Assumptions
- Scheduling Aperiodic Tasks (with soft deadlines)
- Scheduling Aperiodic Tasks (with firm deadlines)

IL2212 – EMBEDDED SOFTWARE

2

2



## Outline

- System Model and Assumptions
- Scheduling Aperiodic Tasks (with soft deadlines)
- Scheduling Aperiodic Tasks (with firm deadlines)

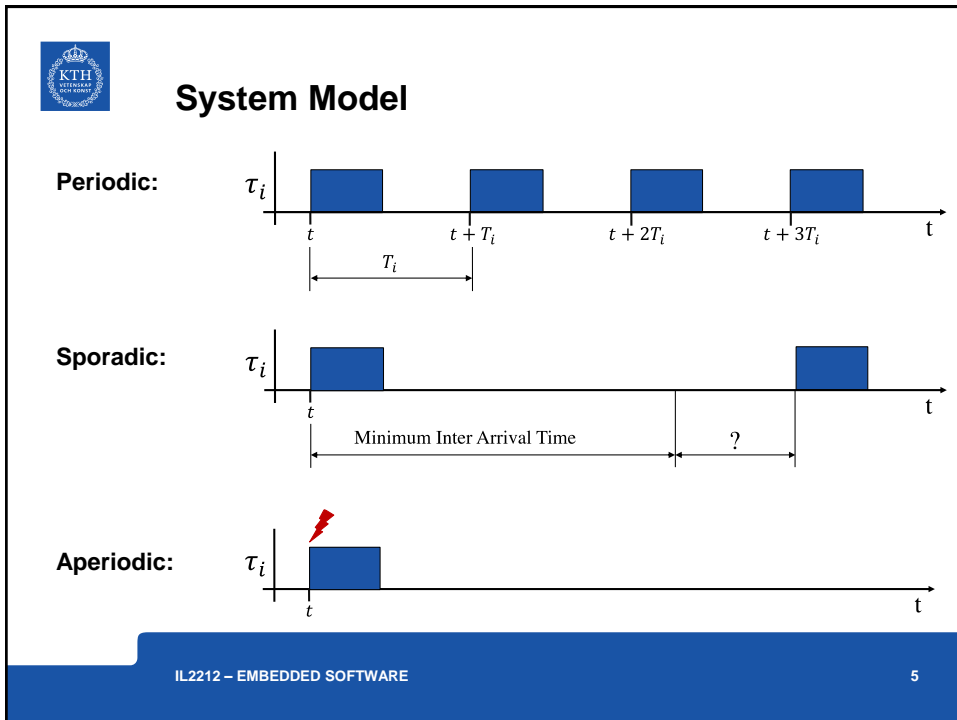
3



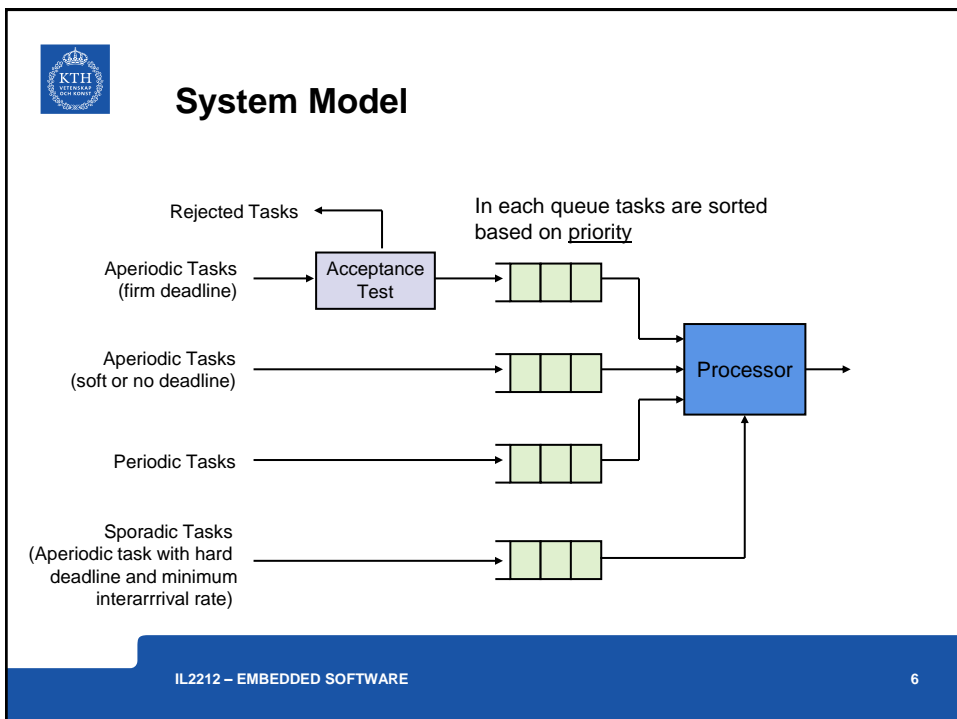
## Assumptions

- Single Processor
- Independent preemptable periodic tasks
- Parameters of all periodic tasks are known
- Periodic tasks meet their deadlines
- Aperiodic and sporadic tasks are independent of each other
- Parameters of aperiodic requests with firm deadlines become known after release

4



5



6



## Scheduling Algorithms

- Aperiodic job requests with soft deadlines
  - are always accepted
  - Scheduler tries to complete aperiodic jobs **as soon as possible**



## Scheduling Algorithms

- Aperiodic job requests with firm deadlines
  - Scheduler decides, if job can be accepted or must be rejected
    - Job is accepted and scheduled, if all other scheduled jobs still meet their deadlines
    - Otherwise, job is rejected
- Sporadic tasks
  - A sporadic task is an aperiodic task with a **hard** deadline
  - Additional information in form of a **minimum interarrival time** as discussed by Buttazzo [But11] is needed!
- Discussion
  - A sporadic job in [Liu00] is an aperiodic job with hard of firm deadline. Lecture uses Buttazzos definition.



## Scheduling Algorithm

- A scheduling algorithm is *correct*, if it only produces correct schedules of the system
- A *correct schedule* is a schedule where periodic tasks and accepted aperiodic jobs with firm deadlines always meet their deadlines



## Optimality of Algorithms

- An aperiodic task scheduling algorithm is optimal if it minimizes either
  - the *response time* of the aperiodic task at the head of the aperiodic task queue, or
  - the *average response time* of all the aperiodic tasks for the given queueing discipline



## Optimality of Algorithms

- A scheduling algorithm for aperiodic tasks with firm deadline is optimal, if it
  - accepts each aperiodic job with firm deadline newly offered to the system and schedules the task to complete in time *if and only if* the new task can be correctly scheduled in time by some means



## Outline

- System Model and Assumptions
- Scheduling Aperiodic Tasks (with soft deadlines)
  - Main principles
  - Background Execution
  - Periodic Server
  - Bandwidth Preserving Server
- Scheduling Aperiodic Tasks (with firm deadlines)



## Scheduling Aperiodic Jobs with Soft Deadline


- In the following several different algorithms to schedule aperiodic jobs are discussed
- Assumption here:
  - No sporadic tasks or aperiodic tasks with firm deadline

13



## Background Execution


- Aperiodic tasks are only scheduled and executed at times, when there is no periodic task ready for execution, i.e. the processor is idle



14



## Background Execution



15



## Background Execution

**Advantage:** simple algorithm

**Disadvantage:** aperiodic jobs are often executed very late

16





## Periodic Server

- A task that behaves more or less like a periodic task and is created to execute aperiodic tasks is called a *periodic server*
- A periodic server is defined partially by execution time  $C_s$  and period  $T_s$
- The parameter  $C_s$  is the execution *capacity* of the server (also called its *budget*)
- The ratio  $U_s = C_s / T_s$  is the *size* of the server


A periodic server can generally be scheduled with the same algorithms as periodic tasks

17



## Periodic Server

- When the server is scheduled it executes aperiodic tasks. While doing that, it *consumes* its budget at the rate of 1 per time unit
- The budget is *exhausted*, when it reaches 0
- A time instant when the budget is replenished (reloaded) is called *replenishment time*



18



## Periodic Server

- A periodic server is *backlogged* whenever the aperiodic task queue is non-empty
- It is *idle* when the queue is empty
- The server is *eligible* for execution only when it is backlogged and has non-zero budget

19



## Polling Server (PS)

- A *polling server* ( $T_s, C_s$ ) is a periodic server
- When executed, it executes an aperiodic task, if the aperiodic task queue is non-empty
- PS suspends execution or is suspended by the scheduler either
  - when it has executed for  $C_s$ , or
  - when there is no aperiodic request pending

20



## Polling Server


- Consumption Rules
  - The budget is **immediately** consumed when the server is not scheduled<sup>1</sup>
- Replenishment Rules
  - The budget is replenished to  $C_s$  at the beginning of each period
- Example: Liu, Figure 7.2b, p.193

<sup>1</sup>If a request arrives just after the server suspended it has to wait for the next period!

21



## Polling Server $\tau_1 = (3,1)$ , $\tau_2 = (10,4)$ , PS = (2.5, 0.5)



22



## Polling Server

- Aperiodic jobs that arrive after the release time of the PS must wait until next polling period
  - Execution budget is **not** preserved
- Simple to prove correctness

23



## Polling Server – Schedulability Analysis

- The schedulability of periodic tasks must be guaranteed in the presence of the polling server.
- From the definition of the PS we can see that
  - In the worst case the server behaves like a periodic task where  $T = T_S$  and  $C = C_S$

$$\sum_{i=1}^n \frac{C_i}{T_i} + \frac{C_S}{T_S} \leq (n+1) \cdot (2^{1/(n+1)} - 1)$$

Polling Server

24



## Bandwidth-Preserving Servers

- A **bandwidth-preserving server** is a periodic server
- Compared to polling server bandwidth preserving servers try to **preserve** their budget when they are not executed
- Additional rules for consumption and replenishment



## Bandwidth-Preserving Servers

- A backlogged bandwidth-preserving server is ready for execution when it has budget
- Scheduler **keeps track** of the consumption of the server budget
- If budget is exhausted server becomes idle



## Bandwidth-Preserving Servers

- Scheduler moves server back to ready queue, when budget is replenished and server is backlogged
- If a new aperiodic job arrives an idle server becomes backlogged and is put into the ready queue when it has budget

27



## Deferrable Server (DS)

- Simplest bandwidth preserving server
- Consumption rule
  - The execution budget of the server is consumed at the rate of one unit per time **whenever the server executes**
- Replenishment rule
  - The execution budget of the server is set to  $C_s$  at multiples of its period
- Server is **not** allowed to cumulate budget from period to period

Introduced to improve the average response time compared to the Polling Server

28





## Schedulability of Deferrable Servers

- Time Demand Analysis can be used to determine whether all tasks remain schedulable in the presence of a deferrable server
- Time Demand Function (if deferrable server has highest priority)

$$w_i(t) = C_i + B_i + C_S + \left\lceil \frac{t - C_S}{T_S} \right\rceil \cdot C_S + \sum_{j=1}^n \left\lceil \frac{t}{T_j} \right\rceil \cdot C_j, \forall 0 \leq t \leq T_i$$

Blocking

Server

IL2212 – EMBEDDED SOFTWARE


31

31



## Deferrable Server Time Demand Analysis

$$\tau_{DS} = (3, 1), \tau_1 = (2, 3.5, 1.5), \tau_2 = (6.5, 0.5)$$



IL2212 – EMBEDDED SOFTWARE

32

32





## Schedulability of Deferrable Servers

- There is no known schedulability utilization that ensures the schedulability of a fixed-priority system in which a deferrable server is scheduled at an arbitrary priority



## Limitations of Deferrable Servers

- Deferrable servers may be scheduled longer than its execution time in a time interval as long as its period
- Lower priority task may be delayed longer by a deferrable server than by a periodic task with same period and execution time

Deferrable server does not behave as a periodic task



## Sporadic Server (SS)

- Bandwidth preserving server
- Improves the average response time of aperiodic tasks without degrading the utilization bound for periodic tasks
- More complex consumption and replenishment rules ensure that each sporadic server with period  $T_s$  and budget  $C_s$  never demands more processor time than the periodic task  $(T_s, C_s)$  in *any* time interval

35



## Sporadic Server – Additional Terms

- The priority of a task that currently executes is  $\pi_{exe}$
- The SS has the associated priority  $\pi_s$
- If  $\pi_{exe} \geq \pi_s$  the SS is **Active**
- If  $\pi_{exe} < \pi_s$  the SS is **Idle**
- The **Replenishment Time** of the SS is **RT**
- The **Replenishment Amount** of the SS is **RA**

36



## Sporadic Server – Replenishment Rules


The Sporadic Server replenishes its capacity based on the following rules:

1. RT is set as soon as SS becomes **active** and  $C_S > 0$ 
  - Example: SS becomes active at time  $t_A$  and  $C_S > 0$ , then  $RT = t_A + T_S$
2. For a certain RT, RA is computed when SS becomes idle or when the capacity is exhausted ( $C_S = 0$ )
  - Example: The SS becomes idle at  $t_I$ . Then RA is set equal to the capacity that has been consumed in the interval  $[t_A, t_I]$

37



## Sporadic Server – Example (Buttazzo P.145)



38



## Sporadic Server – Schedulability Analysis

### Theorem (Sprunt, Sha, Lehoczky): (Buttazzo p. 146)

A periodic task set that is schedulable with a task  $\tau_i$  is also schedulable if  $\tau_i$  is replaced by a Sporadic Server with the same period and execution time.

→ The same schedulability test that was used for the Polling Server can also be used for the Sporadic Server.

$$\sum_{i=1}^n \frac{C_i}{T_i} + \frac{C_S}{T_S} \leq (n+1) \cdot (2^{1/(n+1)} - 1)$$

Sporadic Server



## Fixed-Priority Sporadic Server

- Sporadic server is more complex than polling or deferrable servers due to more complex consumption and replenishment rules
- Main advantage: schedulability easy to demonstrate
- A sporadic server can be treated like a periodic task when we check for schedulability
- System with sporadic server may be schedulable while the corresponding Deferrable Server is not
- More complex sporadic servers exist



## Sporadic Dynamic-Priority Servers

- Sporadic servers can also be used for dynamic-priority (**deadline-driven**) systems, like EDF
- Rules have to be adapted
- Also, here sporadic server can be treated as normal task for schedulability analysis

41



## Other Bandwidth Preserving Servers

- Other bandwidth preserver algorithms are based on **general processor sharing** (GPS) algorithms (deadline-driven algorithms)
- Examples:
  - Constant utilization server
  - Total bandwidth server
  - Weighted round-robin server
- Exact functionality not discussed in course, instead see Buttazzo: 6

42



## Outline

- System Model and Assumptions
- Scheduling Aperiodic Tasks (with soft deadlines)
- Scheduling Aperiodic Tasks (with firm deadlines)

43



## Scheduling Aperiodic Jobs with Firm Deadline

- Aperiodic tasks with firm deadlines
  - Scheduler decides, if a new job request can be accepted or must be rejected
    - Job is accepted and scheduled, if all other scheduled job still meet their deadlines
    - Otherwise, job is rejected
- Aperiodic job with firm deadlines is denoted by  $S_i(a_i, d_i, C_i)$

44



## Acceptance Test in Fixed-Priority System

- Sporadic server can be used to execute sporadic tasks in a fixed-priority system
- The sporadic server  $(T_s, C_s)$  has  $C_s$  units of processor time every  $T_s$  units of time

45



## Acceptance Test in Fixed-Priority Systems

- For each new job  $S(a, d, C)$  of an aperiodic job it must be checked, if
  - new job can be scheduled together with the aperiodic jobs with firm deadline that have deadline before  $d$
  - aperiodic jobs with firm deadline with deadline larger or equal to  $d$  can still be scheduled
- Acceptance test is quite complex, but may still be feasible for many systems

46



## Acceptance Test in Fixed-Priority Systems

- Accepted jobs of aperiodic tasks with firm deadlines are ordered among themselves on **EDF** basis
- For the first aperiodic job with firm deadline  $S_i(t, d_{s,i}, C_{s,i})$  the server has at least  $\left\lfloor \frac{d_{s,i}-t}{T_s} \right\rfloor \cdot C_s$  units of processor time available
- Thus, first job is accepted, if the slack of the task is larger than or equal to 0.

$$\sigma_{s,1}(t) = \left\lfloor \frac{d_{s,1}-t}{T_s} \right\rfloor \cdot C_s - C_{s,1}$$



## Acceptance Test in Fixed-Priority Systems

- When there are already  $n$  accepted aperiodic jobs with a firm deadline in the system, the scheduler computes the slack  $\sigma_{s,i}$  of  $S_i$  according to

$$\sigma_{s,i}(t) = \left\lfloor \frac{d_{s,i}-t}{T_s} \right\rfloor \cdot C_s - C_{s,i} - \sum_{d_{s,k} < d_{s,i}} (C_{s,k} - \xi_{s,k})$$

where  $\xi_{s,k}$  is the execution time of the completed portion of the aperiodic job with firm deadline  $S_k$

( $\xi = \text{xi}$ )





## Acceptance Test in Fixed-Priority Systems

- If the slack  $\sigma_{s,i}$  for the aperiodic job with firm deadline is not less than 0, we have to check, if all accepted jobs can still meet their deadline
- For each aperiodic job with firm deadline  $S_k$  which has an equal or later deadline than  $S_i$ , we have to check, if the slack  $\sigma_{s,k}$  is larger than the execution time of the new aperiodic job  $C_{s,i}$
- The new aperiodic job with firm deadline is only accepted, if this is the case for all accepted aperiodic jobs with firm deadline



## Summary

- Servers can be used for the efficient scheduling of aperiodic jobs with soft and firm deadlines
- Servers have consumption and replenishment rules, which can be arbitrarily complex
- Implementation overhead can be significant



## References

- [1] J. P. Lehoczky, L. Sha, and J. K. Strosnider. Enhanced aperiodic responsiveness in hard real-time environments. In *Proceedings of the IEEE Real-Time Systems Symposium*, December 1987.
- [2] J. K. Strosnider, J. P. Lehoczky, and L. Sha. The deferrable server algorithm for enhancing aperiodic responsiveness in hard-real-time environments. *IEEE Transactions on Computers*, 4(1), January 1995.
- [3] B. Sprunt, L. Sha, and J. Lehoczky. Aperiodic task scheduling for hard-real-time systems. *Journal of Real-Time Systems*, 1(1):27–60, 1989.
- [4] Giorgio C. Buttazzo. *Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and Applications*. Springer, 3rd edition, 2011.
- [5] Jane W. S. Liu. *Real-Time Systems*. Prentice Hall, 2000.