



Resource Access Protocols

IL2212 Embedded Software

Ingo Sander

KTH Royal Institute of Technology
Stockholm, Sweden
`ingo@kth.se`

Outline

- 1 Resource Access Protocols
 - Non-Preemptive Protocol
 - Priority Inheritance Protocol
 - Priority Ceiling Protocol

Outline

- 1 Resource Access Protocols
 - Non-Preemptive Protocol
 - Priority Inheritance Protocol
 - Priority Ceiling Protocol

Resource Access Protocols

- Previous discussion assumed independent tasks
- In practice
 - tasks communicate via communication objects like semaphores, message queues or protected objects with each other
 - \Rightarrow tasks are not independent

Resources

Definition: Resource (Buttazzo, 2011)

- A resource S is any software structure that can be used by a task to advance its execution^a.
- A resource
 - that is dedicated to a particular task is called a **private resource**;
 - that can be used by more tasks is called a **shared resource**;
 - that is protected against concurrent access is called a **exclusive resource**.

^aThe symbol S is used, because the symbol R is already occupied to denote the response time. Also, Buttazzo (2011) uses the semaphore as example for a resource.

- Resources can have several units, e.g. counting semaphore
- Following discussion only deals with exclusive resources with one unit
- Exclusive resources are allocated to tasks on a **non-preemptive basis** and used in a **mutually exclusive manner**

Mutually Exclusive Resource Access

Assumptions

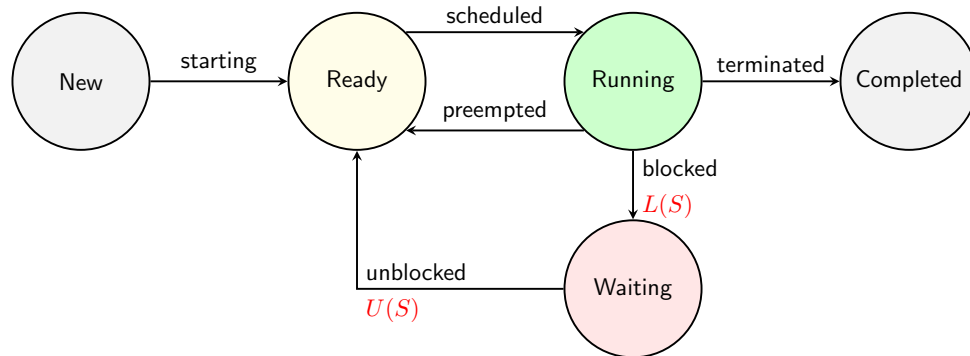
- A lock-based concurrency control mechanism assumed to be used to enforce mutual exclusive access to resources
- When a task wants to use a resource S_i , it executes a command **lock** $L(S_i)$ to request the resource.
- When a task no longer needs the resource S_i , it releases the resource by executing a command **unlock** signal $U(S_i)$.

Correspondence to critical sections protected by a semaphore

A critical section protected by a semaphore S is modelled as an exclusive resource S

- The command **wait**(S) is modelled by a **lock** command $L(S)$
- The command **signal**(S) is modelled by an **unlock** command $U(S)$

Mutually Exclusive Resource Access



- When a lock request $L(S)$ fails, the requesting job is blocked and loses the processor
- An unlock command $U(S)$ frees the resource and might cause that another task is unblocked
- A task stays blocked until the scheduler grants the resources the task is waiting for

Critical Section

- A segment of a task τ_i that begins with a lock $L(S_k)$ and ends with a matching unlock $U(S_k)$ is called a **critical section**, and is denoted by $z_{i,k}$.
- The longest critical section of τ_i to a resource S_k is denoted by $Z_{i,k}$.
- The duration of $Z_{i,k}$ is denoted by $\delta_{i,k}$

...

-- Critical section

lock(S_k)

-- Access to exclusive resource S_k

unlock(S_k)

-- End of critical section

...

Nested Critical Sections

- In **nested critical sections**, resources are always released in last-in-first-out order \Rightarrow **properly nested critical sections**
- $z_{i,h} \subset z_{i,k}$ indicates that the critical section $z_{i,h}$ is entirely contained in $z_{i,k}$.
- A critical section that is not included in other critical sections is called an **outermost critical section**
- Assumption: All nested critical sections are properly nested

Properly Nested Critical Section


```
lock(S_1)
...
lock(S_2)
..
unlock(S_2)
...
unlock(S_1)
```

Not Properly Nested Critical Section

```
lock(S_2)
...
lock(S_1)
..
unlock(S_2)
...
unlock(S_1)
```


Priority Inversion due to Resource Conflicts

- Jobs J_1 , J_2 , J_3 have feasible intervals $(6,14]$, $(2,17]$, $(0,18]$.
- Jobs J_1 , J_2 , J_3 are scheduled using EDF algorithm
- Jobs access critical sections with exclusive resource S_k for the following time units:
 $\delta_{1,k} = 2$, $\delta_{2,k} = 4$, $\delta_{3,k} = 4$.



Priority Inversion due to Resource Conflicts

- Priority inversion occurs, when a low-priority job executes while a ready higher-priority job waits




Timing Anomalies

- Timing Anomalies can occur due to priority inversion
- Assume that critical section of job J_3 is reduced to $\delta_{3,k} = 2.5$ instead of $\delta_{3,k} = 4$.

Timing Anomalies


- Timing Anomalies can occur due to priority inversion
- Assume that critical section of job J_3 is **reduced** to $\delta_{3,k} = 2.5$ instead of $\delta_{3,k} = 4$.



Job J_1 misses its deadline!

Uncontrolled Priority Inversion

- Without suitable protocols, the duration of priority inversion can become unbounded.
- A job J_1 who wants to execute its critical section $z_{1,k}$ can be blocked by a job J_n that holds the resource S_k .
- The duration of $\delta_{i,k}$ can be prolonged by execution of several jobs J_j with a lower priority than the blocked job and a higher priority than the blocking job.
- This phenomenon is called **uncontrolled priority inversion**.

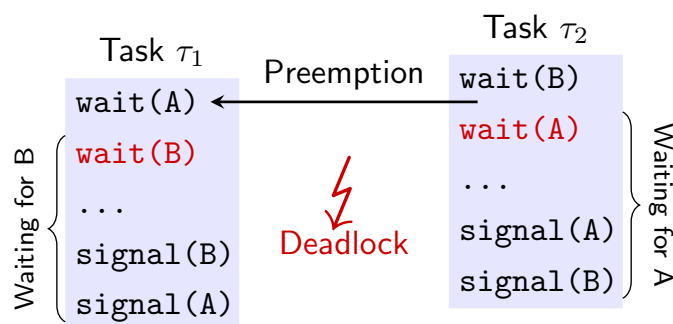


Protocols for Resource Access Control

- Protocols are needed to handle priority inversion in a controlled way and to avoid deadlock

Deadlock

- Deadlock can occur, if tasks block each other from execution




Disciplined approach required to avoid deadlocks

- If a **deadlock** occurs the program cannot proceed!
- Difficult to test for deadlock, because deadlock situation is difficult to create.

Specification of Resource Requirements

- τ_1 needs S_1 for at most 2 time units $\Rightarrow \delta_{1,1} = 2$
- τ_2 needs S_1 for at most 4 time units $\Rightarrow \delta_{2,1} = 4$
- τ_3 needs S_1 for at most 4 time units $\Rightarrow \delta_{3,1} = 4$



Terminology for Resource Access Protocols

- The maximum **blocking time** that task τ_i can experience is denoted by B_i .
- The set of the resources a task τ_i uses is denoted by σ_i .
- The set of resources used by the lower priority task j that can block the task τ_i is denoted by $\sigma_{i,j}$.
- $\gamma_{i,j}$ denotes the set of the longest critical sections of task τ_j that can block task τ_i by accessing a resource S_k .

$$\gamma_{i,j} = \{Z_{j,k} \mid P_j < P_i \wedge S_k \in \sigma_{i,j}\}$$

- γ_i denotes the set of all longest critical sections that can block task τ_i

$$\gamma_i = \bigcup_{j: P_j < P_i} \gamma_{i,j}$$


Non-preemptive Protocol (NPP)

- Simplest resource access control protocol
- All critical sections are scheduled **non-preemptively**, i.e. the **current priority** $p_i(S_k)$ of a task τ_i that starts executing a critical section $z_{i,k}$ is raised to the highest priority level of all tasks.

$$p_i(S_k) = \max_h \{P_h\}$$

- The current priority of the task τ_i is reset to its assigned priority, when the task leaves its critical section.

Non-preemptive Protocol (NPP)



All jobs meet their deadline!

Non-preemptive Protocol (NPP)

- 👍 Simple to implement
- 👍 Uncontrolled priority inversion cannot occur
- 👍 A high-priority job can only be blocked once because of a low-priority job
- 👍 Good protocol when critical sections are short
- 👍 Blocking time can be calculated

$$B_i = \max_{j,k} \{\delta_{j,k} | Z_{j,k} \in \gamma_i\}$$

where

$$\gamma_{i,j} = \{Z_{j,k} | P_j < P_i, k = 1, \dots, m\}$$

Non-preemptive Protocol (NPP)

- 👍 Simple to implement
- 👍 Uncontrolled priority inversion cannot occur
- 👍 A high-priority job can only be blocked once because of a low-priority job
- 👍 Good protocol when critical sections are short
- 👍 Blocking time can be calculated

$$B_i = \max_{j,k} \{\delta_{j,k} | Z_{j,k} \in \gamma_i\}$$

where

$$\gamma_{i,j} = \{Z_{j,k} | P_j < P_i, k = 1, \dots, m\}$$

- 👎 Every task can be blocked by every lower-priority task even if there is no resource conflict between them

Priority Inheritance Protocol

■ Idea

- The blocking jobs J_l **inherits** the priority of the blocked job J_h , when it blocks job J_h
- Avoids that a high priority job J_h is unnecessarily blocked by a low priority job J_l that does use a different resource than the low priority job J_l .

Priority Inheritance Protocol: Definition (Liu, 2000)

■ Definitions

- The priority of a job J_i according to the scheduling algorithm is its assigned priority.
- At any time t , each ready job J_i is scheduled and executes at its current priority $p_i(t)$, which may differ from its assigned priority and vary with time.

■ Rules

- 1 Scheduling Rule
- 2 Allocation Rule
- 3 Priority Inheritance Rule

Scheduling Rule

- Ready jobs are scheduled on the processor preemptively in a priority-driven manner according to their current priorities.
- At its release time t , the current priority $p_i(t)$ of every job J_i is equal to its assigned priority.
- The job remains at this priority except under the condition stated in the priority-inheritance rule.

Allocation Rule


- When a job J_i requests a resource S at time t ,
 - if S is free, S is allocated to J_i until J releases the resource;
 - if S is not free, the request is denied and J_i is blocked.

Priority Inheritance Rule

■ Priority Inheritance


- When the requesting job J_i becomes blocked, the job J_l which blocks J_i inherits the current priority $p_i(t)$ of J_i .
- The job J_l executes at its inherited current priority $p_i(t)$ until it releases the resource S .
- At that time, the priority of J_l returns to its priority $p_l(t')$ when it acquired the resource S .

Priority Inheritance Protocol



All jobs meet their deadline!

Priority Inheritance Protocol



Priority Inheritance Protocol

- 👍 Simple protocol
- 👍 Uncontrolled priority inversion cannot occur

Priority Inheritance Protocol

- 👍 Simple protocol
- 👍 Uncontrolled priority inversion cannot occur
- 👎 Protocol does not minimise the blocking time
 - A high priority job J_h can be blocked by several lower priority jobs, if it runs a nested critical section
 - Difficult to calculate blocking times
- 👎 Protocol does not prevent deadlock. External mechanisms are needed.

Priority Ceiling Protocol

- Extends priority inheritance protocol to prevent deadlocks and aims to further reduce the blocking time
- Key Assumptions
 - Assigned priorities of all jobs are fixed
 - Resources required by all jobs are known a priori before the execution of any job begins

Priority Ceiling Protocol: Definition (Liu, 2000)

■ Definitions

- The **priority ceiling** of any resource S_k is the highest priority of all jobs that require S_k and is denoted $C(S_k)$.
- At any time t , the **current priority ceiling** $C(S^*)(t)$ of the system is
 - equal to the highest priority ceiling of the resources that are in use at the time, if some resources are in use
 - otherwise it is Ω , a non-existing priority that is lower than the priority of any job

■ Rules

- 1 Scheduling Rule
- 2 Allocation Rule
- 3 Priority Inheritance Rule

Scheduling Rule

- At its release time t , the current priority $p_i(t)$ of every job J_i is equal to its assigned priority.
- The job remains at this priority except under the condition stated in the priority-inheritance rule.
- Every job J_i is scheduled preemptively and in a priority-driven manner at its current priority $p_i(t)$.


Allocation Rule

- When a job J_i requests a resource S_k at time t ,
 - 1 if S_k is not free, the request is denied and J_i is blocked.
 - 2 if S_k is free,
 - if J_i 's current priority $p_i(t)$ is higher than the current priority ceiling of the system $C(S^*)(t)$, S_k is allocated to J_i .
 - if J_i 's current priority $p_i(t)$ is not higher than the current ceiling of the system $C(S^*)(t)$, S_k is allocated to J_i only if J_i is the job holding the resource(s) whose priority ceiling is equal to $C(S^*)(t)$; otherwise J_i 's request is denied and J_i becomes blocked.

Priority Inheritance Rule

- Priority Inheritance
 - When J_i becomes blocked, the job J_l which blocks J_i inherits the current priority $p_i(t)$ of J_i .
 - J_l executes at its inherited priority $p_i(t)$ until the time t' where it releases every resource whose priority ceiling is equal to or higher than $p_i(t)$.
 - At that time, the priority of J_l returns to its priority $p_l(t')$ at the time t' when it was granted the resource(s).

Priority Ceiling Protocol



Comparison: Priority Inheritance vs Priority Ceiling Protocol

- Priority-Inheritance protocol is **greedy**: A requesting job will always get access to a resource, if this resource is free
- Priority-Ceiling protocol is **non-greedy**: A requesting job may not get access to a resource, though the resource is free (see additional allocation rule)

Priority Ceiling Protocol: Blocking Types

- Three ways in which a job J_h can be blocked by a low-priority job J_l
 - **Direct Blocking**: A job can be directly blocked by a lower-priority job that owns the requested resource
 - **Priority Inheritance Blocking**: A job J_l can be blocked by a lower-priority job J_l that has inherited the priority of a higher-priority job J_h
 - **Avoidance Blocking**: A job J is blocked by a lower-priority job J_l when J requests a resource S_x that is free at that time, and where J_l currently holds another resource S_y whose priority ceiling is equal to or higher than J 's current priority $p(t)$.

Priority Ceiling Protocol: Important Properties

Freedom from Deadlock (Theorem 8.1 (Liu, 2000))

When resource accesses of a system of preemptive, priority-driven jobs on one processor are controlled by the priority-ceiling protocol, deadlock can never occur.

Priority Ceiling Protocol: Important Properties


! Freedom from Deadlock (Theorem 8.1 (Liu, 2000))

When resource accesses of a system of preemptive, priority-driven jobs on one processor are controlled by the priority-ceiling protocol, deadlock can never occur.

! Blocking Duration (Theorem 8.2 (Liu, 2000))

When resource accesses of preemptive, priority-driven jobs on one processor are controlled by the priority-ceiling protocol, a job can be blocked for at most the duration of one critical section.

Blocking Time Calculation



	Direct Blocking by					Priority Inheritance Blocking by					Priority Ceiling Blocking by				
	J_2	J_3	J_4	J_5	J_6	J_2	J_3	J_4	J_5	J_6	J_2	J_3	J_4	J_5	J_6
J_1		6			2										
J_2	-		5			-	6			2	-	6			2
J_3		-			4		-	5		2		-	5		2
J_4			-					-		4			-		4
J_5				-					-	4				-	

The maximum blocking times are $B_1 = 6$, $B_2 = 6$, $B_3 = 5$, $B_4 = 4$, $B_5 = 4$.

Practical Factors: Blocking Time

An additional term for the **blocking time** B_i can be used in the time demand or response time analysis function. The term B_i reflects the blocking time, which a high priority job can experience when a low priority job executes a non-preemptable section, for instance a critical section protected by a semaphore.

- Time Demand Analysis: $w_i(t) = B_i + C_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{T_k} \right\rceil C_k, \text{ for } 0 < t \leq T_i$
- Response Time Analysis: $R_i = B_i + C_i + \sum_{j=1}^{i-1} \left\lceil \frac{R_i}{T_j} \right\rceil \cdot C_j$

References

- Giorgio C. Buttazzo. [Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and Applications](#). Springer, 3rd edition, 2011.
- Jane W. S. Liu. [Real-Time Systems](#). Prentice Hall, 2000.